

**TRƯỜNG ĐẠI HỌC SƯ PHẠM KỸ THUẬT TP HỒ CHÍ MINH**  
**KHOA CƠ KHÍ CHẾ TẠO MÁY**  
**BỘ MÔN CƠ ĐIỆN TỬ**



**HCMUTE**

**BÀI TẬP VỀ NHÀ**  
**MÔN AI TUẦN 12**

<b>GVHD:</b>	<b>PGS.TS Nguyễn Trường Thịnh</b>
<b>LỚP:</b>	<b>Thứ 7</b>
<b>TIẾT:</b>	<b>3_6</b>
<b>SVTH:</b>	<b>Nguyễn Lê Vũ</b>
<b>MSSV:</b>	<b>19146428</b>

**Tp. Hồ Chí Minh, tháng 05 năm 2022**

# 1. Nhận diện khuôn mặt (dùng ANN)

```
[43] from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout, Conv2D, MaxPooling2D, Flatten
from tensorflow.keras.optimizers import Adam
from sklearn.preprocessing import StandardScaler
from keras.utils import np_utils
from sklearn.utils import shuffle
import cv2
import matplotlib.pyplot as plt
import numpy as np
import pickle
import tensorflow as tf
import math as m
```

```
# Load Data
with open('data.pickle', 'rb') as f:
    (x_train, y_train) = pickle.load(f)
# Reshape Data
x_pre = x_train[130]
x_train = x_train[:194]
y_train = y_train[:194]
x_train = x_train.reshape(x_train.shape[0], -1)
# Preprocessing Data
x_train = x_train.astype('float32')
x_train /= 255
# Encoding Y
y_train = np_utils.to_categorical(y_train, 2)
# Shuffle Data
x_train, y_train = shuffle(x_train, y_train)
```

```
def plot_history(history_fine):
    f1 = history_fine.history['acc']
    val_f1 = history_fine.history['val_acc']
    loss = history_fine.history['loss']
    val_loss = history_fine.history['val_loss']
    plt.figure(figsize=(8, 8))
    plt.subplot(2, 1, 1)
    plt.plot(f1, label='Acc')
    plt.plot(val_f1, label='Validation Acc')
    plt.legend(loc='lower right')
    plt.title('Accuracy')
    plt.subplot(2, 1, 2)
    plt.plot(loss, label='Loss')
    plt.plot(val_loss, label='Validation Loss')
    plt.legend(loc='upper right')
    plt.title('Loss')
    plt.xlabel('epoch')
    plt.show()
```

```
model = Sequential()
model.add(Dense(10, activation='relu', input_shape = (67500,)))
model.add(Dense(10, activation='relu'))
model.add(Dense(10, activation='relu'))
model.add(Dense(10, activation='relu'))
model.add(Dense(10, activation='relu'))
model.add(Dense(10, activation='relu'))
model.add(Dense(10, activation='relu'))
model.add(Dense(10, activation='relu'))
model.add(Dense(10, activation='relu'))
model.add(Dense(2, activation='sigmoid'))



model.compile(loss='binary_crossentropy', optimizer =Adam(), metrics=['acc'])

history = model.fit(x_train, y_train, batch_size = 32, epochs = 50, validation_split = 0.2)
plot_history(history)
```

```
Epoch 1/50
5/5 [=====] - 1s 73ms/step - loss: 0.6682 - acc: 0.5355 - val_loss: 0.6367 - val_acc: 0.4359
Epoch 2/50
5/5 [=====] - 0s 20ms/step - loss: 0.6147 - acc: 0.5355 - val_loss: 0.7268 - val_acc: 0.4359
Epoch 3/50
5/5 [=====] - 0s 16ms/step - loss: 0.5772 - acc: 0.6000 - val_loss: 0.5933 - val_acc: 0.4359
Epoch 4/50
5/5 [=====] - 0s 18ms/step - loss: 0.5107 - acc: 0.6129 - val_loss: 0.4630 - val_acc: 0.6410
Epoch 5/50
5/5 [=====] - 0s 17ms/step - loss: 0.3621 - acc: 0.9290 - val_loss: 0.3275 - val_acc: 1.0000
Epoch 6/50
5/5 [=====] - 0s 18ms/step - loss: 0.2699 - acc: 0.9871 - val_loss: 0.2655 - val_acc: 1.0000
Epoch 7/50
5/5 [=====] - 0s 21ms/step - loss: 0.2016 - acc: 1.0000 - val_loss: 0.1663 - val_acc: 1.0000
Epoch 8/50
5/5 [=====] - 0s 17ms/step - loss: 0.1193 - acc: 1.0000 - val_loss: 0.0891 - val_acc: 1.0000
Epoch 9/50
5/5 [=====] - 0s 17ms/step - loss: 0.0634 - acc: 1.0000 - val_loss: 0.0436 - val_acc: 1.0000
Epoch 10/50
5/5 [=====] - 0s 18ms/step - loss: 0.0334 - acc: 1.0000 - val_loss: 0.0231 - val_acc: 1.0000
Epoch 11/50
5/5 [=====] - 0s 21ms/step - loss: 0.0191 - acc: 1.0000 - val_loss: 0.0133 - val_acc: 1.0000
Epoch 12/50
5/5 [=====] - 0s 21ms/step - loss: 0.0106 - acc: 1.0000 - val_loss: 0.0085 - val_acc: 1.0000
Epoch 13/50
5/5 [=====] - 0s 17ms/step - loss: 0.0068 - acc: 1.0000 - val_loss: 0.0052 - val_acc: 1.0000
Epoch 14/50
5/5 [=====] - 0s 16ms/step - loss: 0.0046 - acc: 1.0000 - val_loss: 0.0039 - val_acc: 1.0000
Epoch 15/50
5/5 [=====] - 0s 16ms/step - loss: 0.0034 - acc: 1.0000 - val_loss: 0.0033 - val_acc: 1.0000
Epoch 16/50
5/5 [=====] - 0s 17ms/step - loss: 0.0027 - acc: 1.0000 - val_loss: 0.0025 - val_acc: 1.0000
Epoch 17/50
5/5 [=====] - 0s 17ms/step - loss: 0.0022 - acc: 1.0000 - val_loss: 0.0021 - val_acc: 1.0000
Epoch 18/50
5/5 [=====] - 0s 19ms/step - loss: 0.0019 - acc: 1.0000 - val_loss: 0.0019 - val_acc: 1.0000
Epoch 19/50
5/5 [=====] - 0s 16ms/step - loss: 0.0017 - acc: 1.0000 - val_loss: 0.0017 - val_acc: 1.0000
Epoch 20/50
5/5 [=====] - 0s 17ms/step - loss: 0.0015 - acc: 1.0000 - val_loss: 0.0015 - val_acc: 1.0000
Epoch 21/50
5/5 [=====] - 0s 16ms/step - loss: 0.0014 - acc: 1.0000 - val_loss: 0.0014 - val_acc: 1.0000
Epoch 22/50
5/5 [=====] - 0s 16ms/step - loss: 0.0012 - acc: 1.0000 - val_loss: 0.0013 - val_acc: 1.0000
Epoch 23/50
5/5 [=====] - 0s 17ms/step - loss: 0.0012 - acc: 1.0000 - val_loss: 0.0011 - val_acc: 1.0000
Epoch 24/50
5/5 [=====] - 0s 17ms/step - loss: 0.0011 - acc: 1.0000 - val_loss: 0.0011 - val_acc: 1.0000
Epoch 25/50
5/5 [=====] - 0s 17ms/step - loss: 0.0010 - acc: 1.0000 - val_loss: 0.0010 - val_acc: 1.0000
Epoch 26/50
5/5 [=====] - 0s 17ms/step - loss: 9.3477e-04 - acc: 1.0000 - val_loss: 9.2967e-04 - val_acc: 1.0000
Epoch 27/50
5/5 [=====] - 0s 17ms/step - loss: 8.7197e-04 - acc: 1.0000 - val_loss: 8.6737e-04 - val_acc: 1.0000
Epoch 28/50
5/5 [=====] - 0s 17ms/step - loss: 8.2622e-04 - acc: 1.0000 - val_loss: 8.2649e-04 - val_acc: 1.0000
Epoch 29/50
5/5 [=====] - 0s 16ms/step - loss: 7.7113e-04 - acc: 1.0000 - val_loss: 7.8008e-04 - val_acc: 1.0000
Epoch 30/50
5/5 [=====] - 0s 18ms/step - loss: 7.3081e-04 - acc: 1.0000 - val_loss: 7.3231e-04 - val_acc: 1.0000
Epoch 31/50
5/5 [=====] - 0s 16ms/step - loss: 6.9722e-04 - acc: 1.0000 - val_loss: 6.8085e-04 - val_acc: 1.0000
Epoch 32/50
5/5 [=====] - 0s 17ms/step - loss: 6.5771e-04 - acc: 1.0000 - val_loss: 6.5709e-04 - val_acc: 1.0000
Epoch 33/50
5/5 [=====] - 0s 16ms/step - loss: 6.2101e-04 - acc: 1.0000 - val_loss: 6.2812e-04 - val_acc: 1.0000
Epoch 34/50
5/5 [=====] - 0s 17ms/step - loss: 5.9010e-04 - acc: 1.0000 - val_loss: 5.9735e-04 - val_acc: 1.0000
Epoch 35/50
5/5 [=====] - 0s 17ms/step - loss: 5.6276e-04 - acc: 1.0000 - val_loss: 5.6177e-04 - val_acc: 1.0000
Epoch 36/50
5/5 [=====] - 0s 17ms/step - loss: 5.3553e-04 - acc: 1.0000 - val_loss: 5.3193e-04 - val_acc: 1.0000
```



Epoch 36/50  
 5/5 [=====] - 0s 17ms/step - loss: 5.3553e-04 - acc: 1.0000 - val\_loss: 5.3193e-04 - val\_acc: 1.0000  
 Epoch 37/50  
 5/5 [=====] - 0s 17ms/step - loss: 5.1273e-04 - acc: 1.0000 - val\_loss: 5.0614e-04 - val\_acc: 1.0000  
 Epoch 38/50  
 5/5 [=====] - 0s 18ms/step - loss: 4.8696e-04 - acc: 1.0000 - val\_loss: 4.8717e-04 - val\_acc: 1.0000  
 Epoch 39/50  
 5/5 [=====] - 0s 17ms/step - loss: 4.6998e-04 - acc: 1.0000 - val\_loss: 4.7692e-04 - val\_acc: 1.0000  
 Epoch 40/50  
 5/5 [=====] - 0s 17ms/step - loss: 4.4780e-04 - acc: 1.0000 - val\_loss: 4.5340e-04 - val\_acc: 1.0000  
 Epoch 41/50  
 5/5 [=====] - 0s 21ms/step - loss: 4.2919e-04 - acc: 1.0000 - val\_loss: 4.3321e-04 - val\_acc: 1.0000  
 Epoch 42/50  
 5/5 [=====] - 0s 17ms/step - loss: 4.1075e-04 - acc: 1.0000 - val\_loss: 4.0877e-04 - val\_acc: 1.0000  
 Epoch 43/50  
 5/5 [=====] - 0s 17ms/step - loss: 3.9347e-04 - acc: 1.0000 - val\_loss: 3.9020e-04 - val\_acc: 1.0000  
 Epoch 44/50  
 5/5 [=====] - 0s 18ms/step - loss: 3.7969e-04 - acc: 1.0000 - val\_loss: 3.7981e-04 - val\_acc: 1.0000  
 Epoch 45/50  
 5/5 [=====] - 0s 17ms/step - loss: 3.6347e-04 - acc: 1.0000 - val\_loss: 3.6680e-04 - val\_acc: 1.0000  
 Epoch 46/50  
 5/5 [=====] - 0s 17ms/step - loss: 3.4942e-04 - acc: 1.0000 - val\_loss: 3.5173e-04 - val\_acc: 1.0000  
 Epoch 47/50  
 5/5 [=====] - 0s 18ms/step - loss: 3.3845e-04 - acc: 1.0000 - val\_loss: 3.4135e-04 - val\_acc: 1.0000  
 Epoch 48/50  
 5/5 [=====] - 0s 17ms/step - loss: 3.2418e-04 - acc: 1.0000 - val\_loss: 3.2750e-04 - val\_acc: 1.0000  
 Epoch 49/50  
 5/5 [=====] - 0s 16ms/step - loss: 3.1209e-04 - acc: 1.0000 - val\_loss: 3.1307e-04 - val\_acc: 1.0000  
 Epoch 50/50  
 5/5 [=====] - 0s 18ms/step - loss: 3.0102e-04 - acc: 1.0000 - val\_loss: 3.0144e-04 - val\_acc: 1.0000




✓  
1s



```
# Load Test Image
plt.imshow(cv2.cvtColor(x_pre, cv2.COLOR_BGR2RGB))
print(x_pre.shape)
img = x_pre.reshape(1,-1)
img = img.astype('float32')
img /= 255
```

↳ (150, 150, 3)




✓  
2s



```
plt.title("Model dự đoán (1: Vũ, 0: Người lạ): " + str(np.argmax(model.predict(img))))
plt.imshow(cv2.cvtColor(x_pre, cv2.COLOR_BGR2RGB), cmap=plt.get_cmap('gray'))
```

↳ <matplotlib.image.AxesImage at 0x7f3da1ebc650>

Model dự đoán (1: Vũ, 0: Người lạ): 1



## 2. Robot 2 bậc


```
[2] from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout, Conv2D, MaxPooling2D, Flatten
from tensorflow.keras.optimizers import Adam
from sklearn.preprocessing import StandardScaler
from keras.utils import np_utils
from sklearn.utils import shuffle
import cv2
import matplotlib.pyplot as plt
import numpy as np
import pickle
import tensorflow as tf
import math as m
```

```
def plot_reg_history(history_fine):
    loss = history_fine.history['loss']
    val_loss = history_fine.history['val_loss']
    plt.subplot(2, 1, 2)
    plt.plot(loss, label='Loss')
    plt.plot(val_loss, label='Validation Loss')
    plt.legend(loc='upper right')
    plt.title('Loss')
    plt.xlabel('epoch')
    plt.show()

# Define Variables
l1 = 40
l2 = 50
x_train = []
y_train = []
# Create Data
for t1 in np.linspace(-(2 * np.pi), 2 * np.pi, 500):
    for t2 in np.linspace(-(2 * np.pi), 2 * np.pi, 500):
        x = l1*m.cos(t1) + l2*m.cos(t1+t2)
        y = l1*m.sin(t1) + l2*m.sin(t1+t2)
        x_train.append(np.array([x,y]))
        y_train.append(np.array([t1,t2]))
# Convert to array
scaler = StandardScaler()
x_train = np.array(scaler.fit_transform(x_train))
y_train = np.array(y_train)
# Shuffle
x_train, y_train = shuffle(x_train, y_train)
```

```
[11] model = Sequential()
model.add(Dense(256, activation='relu', input_shape = (2,)))
model.add(Dense(256, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(2, activation='linear'))
model.compile(loss='mae', optimizer=tf.optimizers.Adam(learning_rate=0.0001))
history = model.fit(x_train, y_train, batch_size = 512, epochs = 10, validation_split = 0.2)
plot_reg_history(history)
```

Epoch 1/10  
 391/391 [=====] - 3s 6ms/step - loss: 3.1495 - val\_loss: 3.1416  
 Epoch 2/10  
 391/391 [=====] - 2s 4ms/step - loss: 3.1495 - val\_loss: 3.1416  
 Epoch 3/10  
 391/391 [=====] - 2s 4ms/step - loss: 3.1495 - val\_loss: 3.1416  
 Epoch 4/10  
 391/391 [=====] - 2s 4ms/step - loss: 3.1495 - val\_loss: 3.1416  
 Epoch 5/10  
 391/391 [=====] - 2s 4ms/step - loss: 3.1495 - val\_loss: 3.1416  
 Epoch 6/10  
 391/391 [=====] - 2s 4ms/step - loss: 3.1495 - val\_loss: 3.1416  
 Epoch 7/10  
 391/391 [=====] - 2s 4ms/step - loss: 3.1495 - val\_loss: 3.1416  
 Epoch 8/10  
 391/391 [=====] - 2s 4ms/step - loss: 3.1495 - val\_loss: 3.1416  
 Epoch 9/10  
 391/391 [=====] - 2s 4ms/step - loss: 3.1495 - val\_loss: 3.1416  
 Epoch 10/10  
 391/391 [=====] - 2s 4ms/step - loss: 3.1495 - val\_loss: 3.1416



```
test = scaler.transform(np.array([[90,0]]))
t1 = model.predict(test)[0][0]
t2 = model.predict(test)[0][1]

x = l1*m.cos(t1) + l2*m.cos(t2+t1)
y = l1*m.sin(t1) + l2*m.sin(t2+t1)

print("Model dự đoán với giá trị đầu vào x = 90 và y = 0 là t1 = " + str(t1) + " t2 = " + str(t2))
print("Kiểm tra: ")
print("Với giá trị t1 và t2 dự đoán ta tính lại x = " + str(x) + " y = " + str(y))
```

Model dự đoán với giá trị đầu vào x = 90 và y = 0 là t1 = 0.00045843548 t2 = 0.0011386004  
 Kiểm tra:  
 Với giá trị t1 và t2 dự đoán ta tính lại x = 89.99993203365955 y = 0.09818917989102906

### 3. Robot 3 bậc

```
[7] from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout, Conv2D, MaxPooling2D, Flatten
from tensorflow.keras.optimizers import Adam
from sklearn.preprocessing import StandardScaler
from keras.utils import np_utils
from sklearn.utils import shuffle
import cv2
import matplotlib.pyplot as plt
import numpy as np
import pickle
import tensorflow as tf
import math as m
```

```
# Define Variables
l1 = 40
l2 = 50
l3 = 20
x_train = []
y_train = []
# Create Data
for t1 in np.linspace(-(2 * np.pi), 2 * np.pi, 100):
    for t2 in np.linspace(-(2 * np.pi), 2 * np.pi, 100):
        for t3 in np.linspace(-(2 * np.pi), 2 * np.pi, 100):
            x = l1*m.cos(t1) + l2*m.cos(t1+t2) + l3*m.cos(t1+t2+t3)
            y = l1*m.sin(t1) + l2*m.sin(t1+t2) + l3*m.sin(t1+t2+t3)
            beta = (t1 + t2 + t3)*180/3.14
            x_train.append(np.array([x,y,beta]))
            y_train.append(np.array([t1,t2,t3]))
# Convert to array
scaler = StandardScaler()
x_train = np.array(scaler.fit_transform(x_train))
y_train = np.array(y_train)
# Shuffle
x_train, y_train = shuffle(x_train, y_train)
```

```
[9] def plot_reg_history(history_fine):
    loss = history_fine.history['loss']
    val_loss = history_fine.history['val_loss']
    plt.subplot(2, 1, 2)
    plt.plot(loss, label='Loss')
    plt.plot(val_loss, label='Validation Loss')
    plt.legend(loc='upper right')
    plt.title('Loss')
    plt.xlabel('epoch')
    plt.show()
```



2m

```


model = Sequential()
model.add(Dense(256, activation='relu', input_shape = (3,)))
model.add(Dense(256, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(3, activation='linear'))
model.compile(loss='mae', optimizer = tf.optimizers.Adam(learning_rate=0.0001))
history = model.fit(x_train, y_train, batch_size = 512, epochs = 10, validation_split = 0.2)
plot_reg_history(history)

```

```

Epoch 1/10
1563/1563 [=====] - 14s 9ms/step - loss: 2.3900 - val_loss: 2.3321
Epoch 2/10
1563/1563 [=====] - 15s 9ms/step - loss: 2.3216 - val_loss: 2.2974
Epoch 3/10
1563/1563 [=====] - 13s 9ms/step - loss: 2.2830 - val_loss: 2.2620
Epoch 4/10
1563/1563 [=====] - 13s 9ms/step - loss: 2.2507 - val_loss: 2.2450
Epoch 5/10
1563/1563 [=====] - 13s 8ms/step - loss: 2.2351 - val_loss: 2.2291
Epoch 6/10
1563/1563 [=====] - 13s 8ms/step - loss: 2.2264 - val_loss: 2.2232
Epoch 7/10
1563/1563 [=====] - 14s 9ms/step - loss: 2.2212 - val_loss: 2.2190
Epoch 8/10
1563/1563 [=====] - 15s 9ms/step - loss: 2.2172 - val_loss: 2.2124
Epoch 9/10
1563/1563 [=====] - 14s 9ms/step - loss: 2.2144 - val_loss: 2.2095
Epoch 10/10
1563/1563 [=====] - 14s 9ms/step - loss: 2.2110 - val_loss: 2.2150

```



0s

```

test = scaler.transform(np.array([[60,0,45]]))
t1 = model.predict(test)[0][0]
t2 = model.predict(test)[0][1]
t3 = model.predict(test)[0][2]
x = 11*m.cos(t1) + 12*m.cos(t1+t2) + 13*m.cos(t1+t2+t3)
y = 11*m.sin(t1) + 12*m.sin(t1+t2) + 13*m.sin(t1+t2+t3)
beta = (t1 + t2 + t3)*180/3.14
print("Model dự đoán với giá trị đầu vào x = 90, y = 0 và beta = 45 là t1 = " + str(t1) + " t2 = " + str(t2) + " t3 = " + str(t3))
print("Kiểm tra: ")
print("Với giá trị t1 và t2 dự đoán ta tính lại x = " + str(x) + " y = " + str(y) + " beta = " + str(beta))

```

## 4. Nhận diện khuôn mặt (dùng CNN)

```
[27] from keras.models import Sequential
      from keras.layers import Dense, Activation, Dropout, Conv2D, MaxPooling2D, Flatten
      from tensorflow.keras.optimizers import Adam
      from sklearn.preprocessing import StandardScaler
      from keras.utils import np_utils
      from sklearn.utils import shuffle
      import cv2
      import matplotlib.pyplot as plt
      import numpy as np
      import pickle
      import tensorflow as tf
      import math as m
```

```
[28] with open('data.pickle', 'rb') as f:
      (x_train, y_train) = pickle.load(f)
      x_pre_1 = x_train[14]
      x_pre_2 = x_train[196]
      x_pre_3 = x_train[220]
      x_train = x_train.astype('float32')
      x_train /= 255
      y_train = np_utils.to_categorical(y_train, 3)
      x_train, y_train = shuffle(x_train, y_train)
```

```
def plot_history(history_fine):
    f1 = history_fine.history['acc']
    val_f1 = history_fine.history['val_acc']
    loss = history_fine.history['loss']
    val_loss = history_fine.history['val_loss']
    plt.figure(figsize=(8, 8))
    plt.subplot(2, 1, 1)
    plt.plot(f1, label='Acc')
    plt.plot(val_f1, label='Validation Acc')
    plt.legend(loc='lower right')
    plt.title('Accuracy')
    plt.subplot(2, 1, 2)
    plt.plot(loss, label='Loss')
    plt.plot(val_loss, label='Validation Loss')
    plt.legend(loc='upper right')
    plt.title('Loss')
    plt.xlabel('epoch')
    plt.show()
```

```

model = Sequential()
model.add(Conv2D(32, (3,3), activation='relu',kernel_initializer='he_uniform', padding = 'same', input_shape = (150,150,3)))
model.add(Conv2D(32, (3,3), activation='relu',kernel_initializer='he_uniform', padding = 'same'))
model.add(MaxPooling2D(2,2))
model.add(Conv2D(64, (3,3), activation='relu',kernel_initializer='he_uniform', padding = 'same'))
model.add(Conv2D(64, (3,3), activation='relu',kernel_initializer='he_uniform', padding = 'same'))
model.add(MaxPooling2D(2,2))
model.add(Conv2D(128, (3,3), activation='relu',kernel_initializer='he_uniform', padding = 'same'))
model.add(Conv2D(128, (3,3), activation='relu',kernel_initializer='he_uniform', padding = 'same'))
model.add(MaxPooling2D(2,2))
model.add(Flatten())
model.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))
model.add(Dense(3, activation='softmax'))
model.summary()

```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_12 (Conv2D)	(None, 150, 150, 32)	896
conv2d_13 (Conv2D)	(None, 150, 150, 32)	9248
max_pooling2d_6 (MaxPooling 2D)	(None, 75, 75, 32)	0
conv2d_14 (Conv2D)	(None, 75, 75, 64)	18496
conv2d_15 (Conv2D)	(None, 75, 75, 64)	36928
max_pooling2d_7 (MaxPooling 2D)	(None, 37, 37, 64)	0
conv2d_16 (Conv2D)	(None, 37, 37, 128)	73856
conv2d_17 (Conv2D)	(None, 37, 37, 128)	147584
max_pooling2d_8 (MaxPooling 2D)	(None, 18, 18, 128)	0
flatten_2 (Flatten)	(None, 41472)	0
dense_4 (Dense)	(None, 128)	5308544
dense_5 (Dense)	(None, 3)	387
=====		
Total params: 5,595,939		
Trainable params: 5,595,939		
Non-trainable params: 0		

```

[20] opt = Adam(lr = 0.001)
model.compile(optimizer = opt, loss = 'categorical_crossentropy', metrics = ['acc'])
his = model.fit(x_train, y_train, epochs = 10, batch_size = 64, validation_split = 0.2)

```

```

Epoch 1/10
/usr/local/lib/python3.7/dist-packages/keras/optimizer_v2/adam.py:105: UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
  super(Adam, self).__init__(name, **kwargs)
4/4 [=====] - 1s 196ms/step - loss: 9.5859e-07 - acc: 1.0000 - val_loss: 0.0000e+00 - val_acc: 1.0000
Epoch 2/10
4/4 [=====] - 1s 132ms/step - loss: 0.0000e+00 - acc: 1.0000 - val_loss: 0.0000e+00 - val_acc: 1.0000
Epoch 3/10
4/4 [=====] - 1s 135ms/step - loss: 0.0000e+00 - acc: 1.0000 - val_loss: 0.0000e+00 - val_acc: 1.0000
Epoch 4/10
4/4 [=====] - 1s 131ms/step - loss: 0.0000e+00 - acc: 1.0000 - val_loss: 0.0000e+00 - val_acc: 1.0000
Epoch 5/10
4/4 [=====] - 1s 134ms/step - loss: 2.9802e-09 - acc: 1.0000 - val_loss: 0.0000e+00 - val_acc: 1.0000
Epoch 6/10
4/4 [=====] - 1s 134ms/step - loss: 0.0000e+00 - acc: 1.0000 - val_loss: 0.0000e+00 - val_acc: 1.0000
Epoch 7/10
4/4 [=====] - 1s 138ms/step - loss: 0.0000e+00 - acc: 1.0000 - val_loss: 0.0000e+00 - val_acc: 1.0000
Epoch 8/10
4/4 [=====] - 1s 135ms/step - loss: 0.0000e+00 - acc: 1.0000 - val_loss: 0.0000e+00 - val_acc: 1.0000
Epoch 9/10
4/4 [=====] - 1s 159ms/step - loss: 0.0000e+00 - acc: 1.0000 - val_loss: 0.0000e+00 - val_acc: 1.0000
Epoch 10/10
4/4 [=====] - 1s 137ms/step - loss: 0.0000e+00 - acc: 1.0000 - val_loss: 0.0000e+00 - val_acc: 1.0000


```

```

[33] label = ['Duc', 'Vu', 'Phat']
plt.title("Du doan nguoi trong anh la: " + label[np.argmax(model.predict(x_pre_1.reshape(1,150,150,3)))])
plt.imshow(cv2.cvtColor(x_pre_1, cv2.COLOR_BGR2RGB), cmap=plt.get_cmap('gray'))

```

<matplotlib.image.AxesImage at 0x7fe56dd1f210>




```

[34] plt.title("Du doan nguoi trong anh la: " + label[np.argmax(model.predict(x_pre_2.reshape(1,150,150,3)))])
plt.imshow(cv2.cvtColor(x_pre_2, cv2.COLOR_BGR2RGB), cmap=plt.get_cmap('gray'))

```

<matplotlib.image.AxesImage at 0x7fe56dcf05d0>



```

[35] plt.title("Du doan nguoi trong anh la: " + label[np.argmax(model.predict(x_pre_3.reshape(1,150,150,3)))])
plt.imshow(cv2.cvtColor(x_pre_3, cv2.COLOR_BGR2RGB), cmap=plt.get_cmap('gray'))

```

<matplotlib.image.AxesImage at 0x7fe56dc6b0d0>

